

Machine Learning
Developer Day

January 24, 2019
12:00-5:00pm

C L Kim

Verizon Media (FKA Oath Inc./Aol/Yahoo)

**USING ML IN AN ANDROID DEMO APP
– THE GOOD, THE BAD, AND THE UGLY**

OVERVIEW

- ▶ Code for the demo app on github github.com/clkim/NumberDetector
 - ▶ Demo app simply tries to classify hand-drawn digits
- ▶ This talk highlights some of the Good, the Bad, and the Ugly
 - ▶ Journey using a trained ML model to build an Android demo app
 - ▶ Some code (Kotlin), and some details of ML model needed to develop app
 - ▶ Time constraint: only covering code related to working with the ML model
- ▶ We do not cover in any depth:
 - ▶ Building/Training a ML model (Google TensorFlow Lite)
 - ▶ APIs: Android; Firebase (Google's mobile dev platform: iOS, Android, Web)

ACKNOWLEDGEMENTS AND CREDITS

- ▶ Mark Allison: *ML for Android Developers* – Part 1, 2, 3
 - ▶ The inspiration and material for this talk, as well as the source code for the demo app, are largely from the above blog series blog.stylingandroid.com/ml-for-android-developers-part-1-2/
- ▶ Tianxing Li: github repo *MNIST with TensorFlow Lite on Android*
 - ▶ The TensorFlow model trained on the MNIST dataset, and converted to TensorFlow Lite github.com/nex3z/tflite-mnist-android
 - ▶ Uses *tf.nn* instead of *tf.keras.layers* module used by TensorFlow “official” MNIST model; seems equivalent

SOME KOTLIN (1 / 3)

- ▶ Kotlin is JVM language that Google seems to promote for Android development developer.android.com/kotlin/
- ▶ “Pascal syntax” so the type comes after name and colon

```
fun classify(  
    bitmap: Bitmap,  
    success: (Int, Float, Long) -> Unit  
)
```



SOME KOTLIN (2 / 3)

► Extension Functions with Receiver Type (inspired by C#)

```
private val imagePixels = IntArray(imageSize)
...
```

```
private fun Bitmap.toVector(): Array<Array<Array<FloatArray>>> {
    getPixels(imagePixels, 0, width, 0, 0, width, height)
    return Array(1) {
        Array(imageHeight) { y ->
            Array(imageWidth) { x ->
                floatArrayOf(imagePixels[x + (y * imageWidth)]
                    .convertToGreyScale()
            )
        }
    }
}
```

SOME KOTLIN (3 / 3)

- ▶ Safe Calls (null safety) with safe call operator ?.

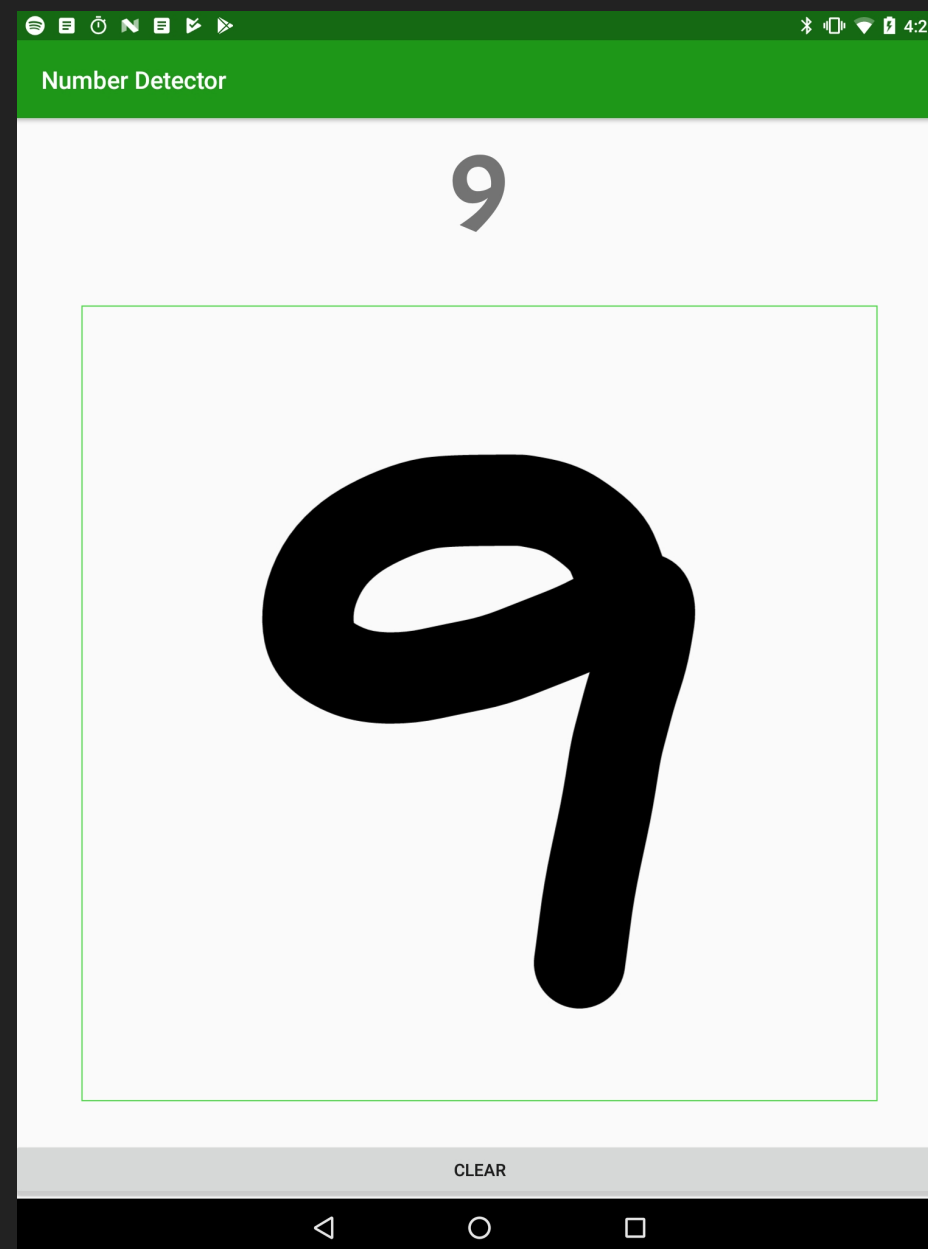
```
interpreter.run(inputs, options)
    .addOnSuccessListener { outputs ->
        outputs.map().entries.maxBy { it.value }
            ?.also { maxEntry ->
                success(maxEntry.key, maxEntry.value,
                    System.currentTimeMillis() - start)
            }
    }
```

- ▶ `maxEntry?.also({<lambda>})`
 - ▶ `also({<lambda>})` is not called if `maxEntry == null`, no NPE;
 - ▶ `also({<lambda>})` is called otherwise

SOME ML FRAMEWORKS/LIBRARIES

- ▶ Mobile focused:
 - ▶ Google TensorFlow Lite, Apple Core ML, Caffe2(Facebook), Xmartlabs Bender, Quantized-CNN
- ▶ Server focused:
 - ▶ TensorFlow, Amazon AML, Microsoft Cognitive Toolkit, Apache MXNet, Apache Spark MLlib
- ▶ Source: <http://hopinfirst.com/top-10-machine-learning-frameworks-mobile-apps/>

SIMPLE ANDROID DEMO APP



Hand-draw a digit, multiple motions ok

Acknowledgement:

Mark Allison's blog posts and sample code is the inspiration and source for this talk

<https://blog.stylingandroid.com/ml-for-android-developers-part-1-2/>

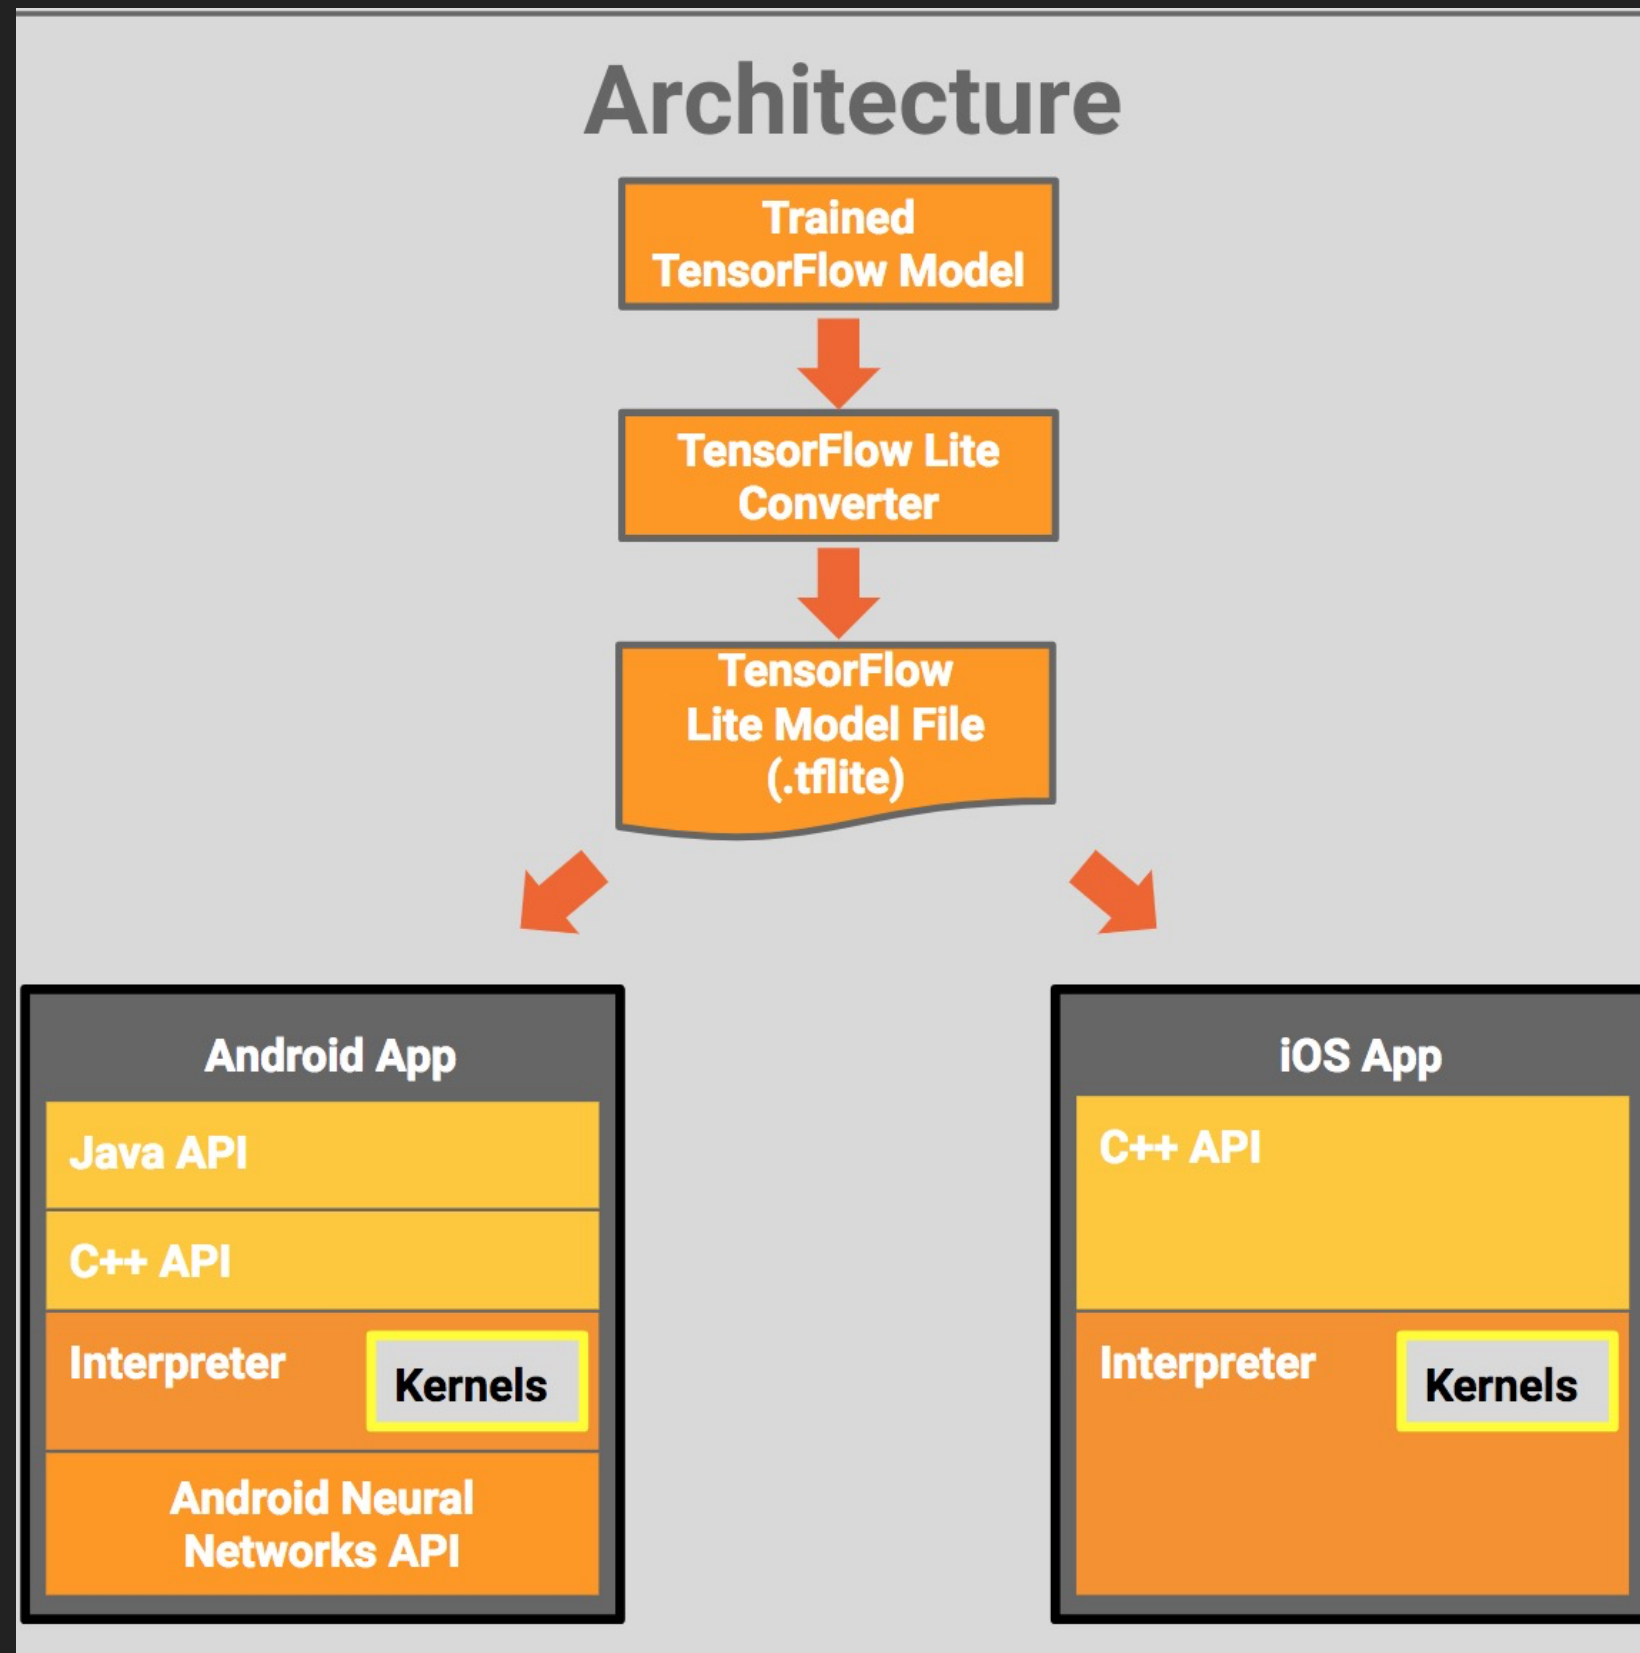
OVERVIEW – ANDROID DEMO APP USING MACHINE LEARNING



TENSORFLOW LITE

- ▶ TensorFlow for mobile/embedded devices
- ▶ Low latency, small binary
- ▶ Android iOS RaspberryPi
- ▶ On-device interpreter
- ▶ Supports hw acceleration
- ▶ Ref: TensorFlow Lite > GUIDE tab (TFLG)

www.tensorflow.org/lite/overview



USE TF LITE MODEL DIRECTLY, OR HOSTED BY ML KIT (BETA)

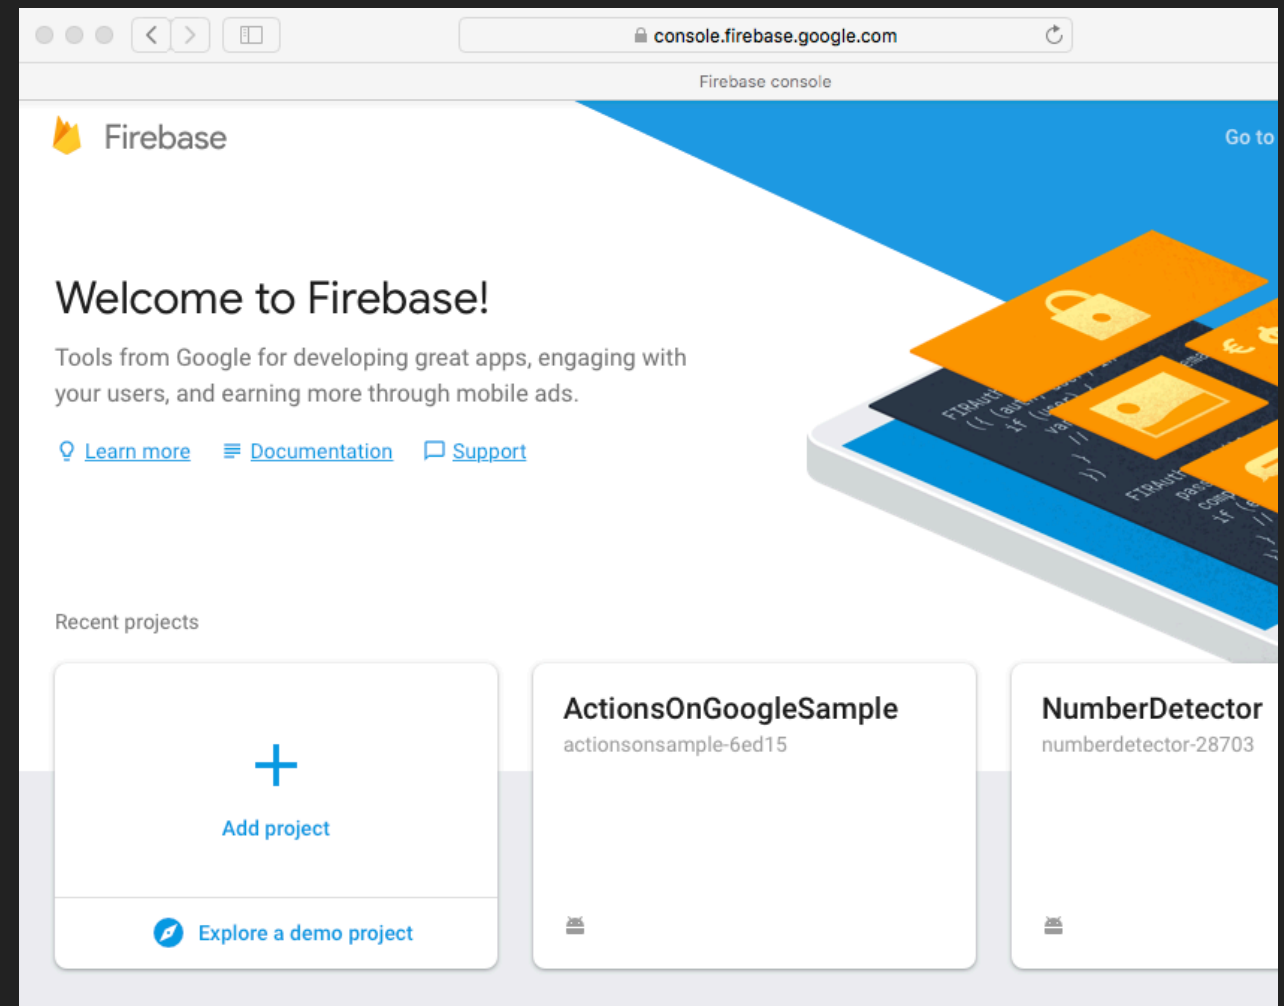
- ▶ TFLG > Android demo app: camera demo app; classifies images continuously captured by camera
 - ▶ The build process downloads the trained model, eg. quantized MobileNet TensorFlow Lite, and bundles it into the app binary
- ▶ We use ML Kit (beta release) for Firebase to host trained model
 - ▶ Decouple the model from the app
 - ▶ **The Good:** An upgraded model can be released at any time, without an app update to Google Play (or Apple App Store)

GETTING A TF LITE MODEL

- ▶ TFLG > Developer guide: good overview of
 - ▶ Training a pre-trained or custom model
 - ▶ Converting to TensorFlow Lite model
- ▶ Our Demo App takes a short-cut: it does use the MNIST dataset of hand-drawn numeric digits; but instead of starting from scratch, we found a trained model already converted to TensorFlow Lite, e.g. *mnist.tflite* (by Tianxing Li, github.com/nex3z/tflite-mnist-android)

ML KIT

- ▶ First, create a Firebase account: login at console.firebase.google.com
- ▶ > Add project, e.g. NumberDetector
- ▶ > Download google-services.json
> follow paged instructions to update the app project in Android Studio

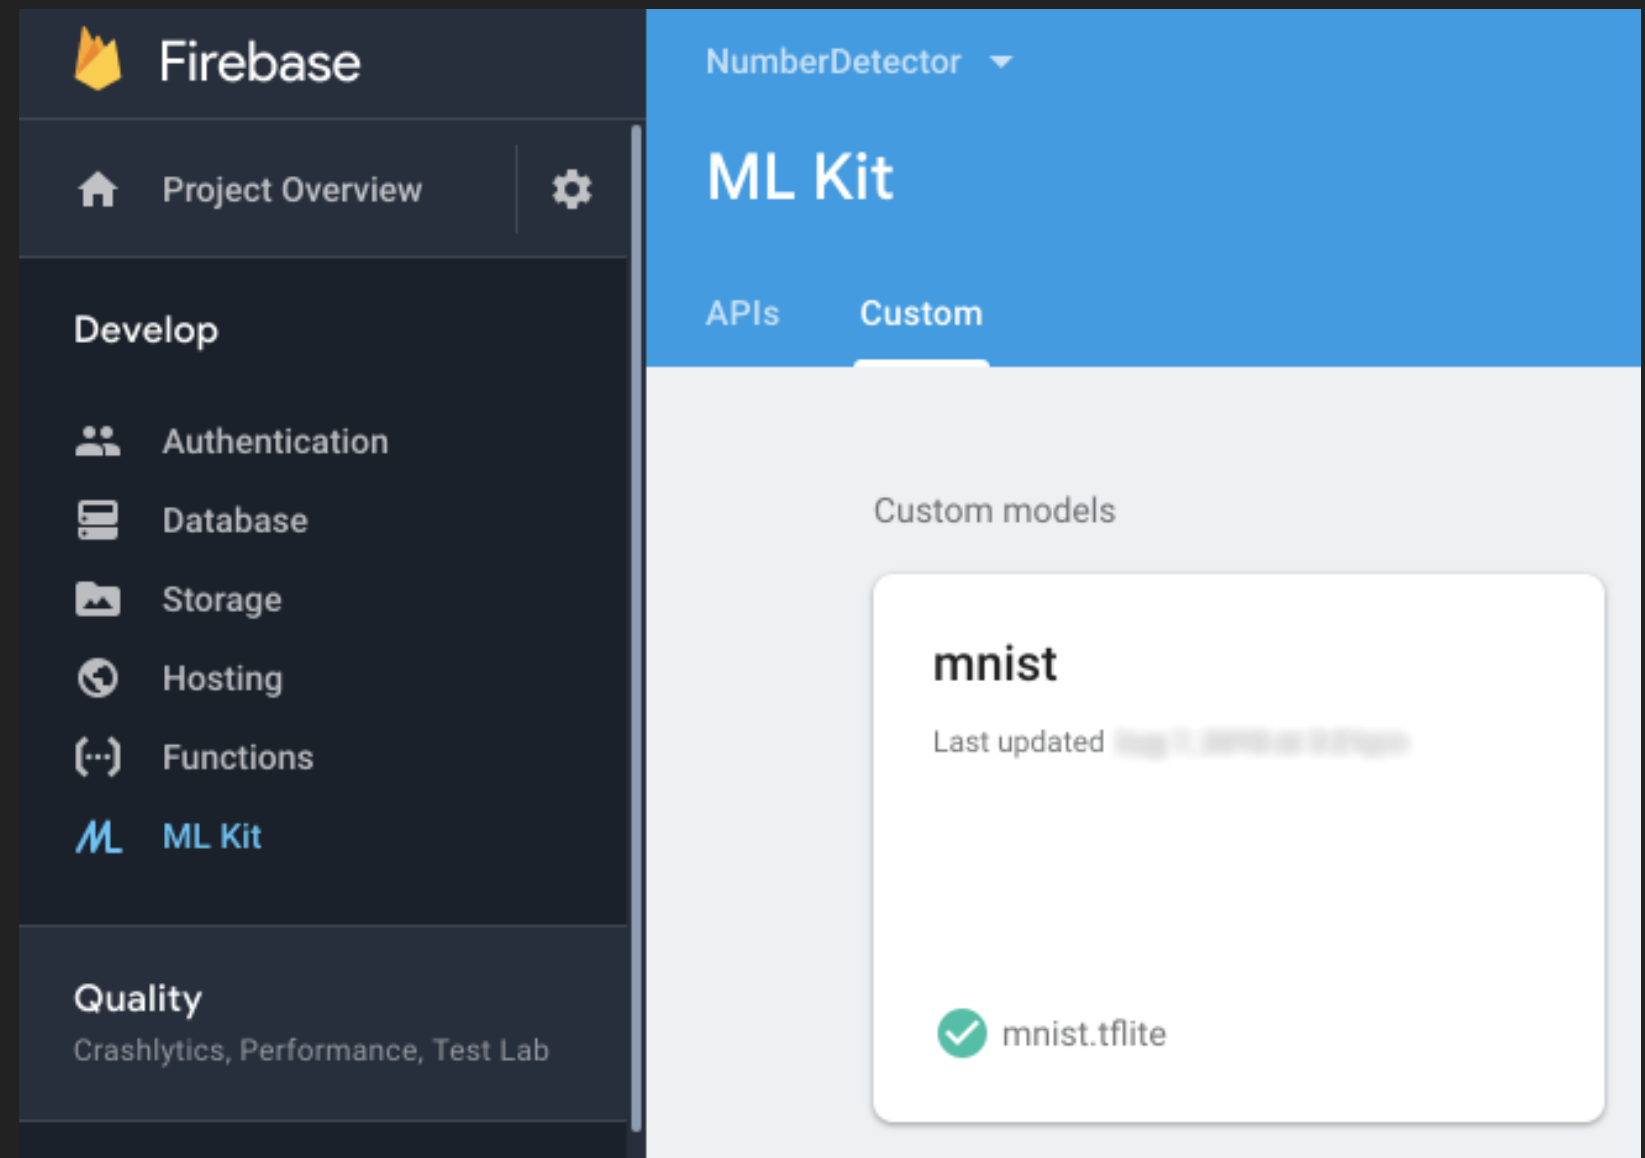


ML KIT (2)

- ▶ Last step: connect the app to Firebase by running the app from IDE > Run | Run app; but Catch-22 is: what if app is not ready/debugged?
- ▶ A hack, try: use Firebase Assistant to pick any (benign) Firebase feature in order to connect the app, so from IDE > Tool | Firebase > e.g. Test Lab > Run Firebase Test Lab...

ML KIT (3)

- ▶ Upload the *mnist.tflite* file to Firebase
- ▶ Firebase console > ML Kit > Custom tab



Firebase console

MAGIC SAUCE – THE INPUT AND OUTPUT FORMATS

- ▶ A TF Lite model expects as input, and produces as output
 - ▶ One or more multidimensional arrays
 - ▶ Values of type: *byte*, *int*, *long*, or *float*
- ▶ **The Ugly**: How to know what are their “shape” and type?
 - ▶ 1. Ask the Machine Learning team which trained the model and provided the *.tflite* file
 - ▶ 2. Find out from the *.tflite* file

GRAPH VISUALIZATION TOOL

▶ There is tool to visualize TensorFlow Lite models...

▶ if you can find/build it :)

1. Install Bazel build tool, see: docs.bazel.build/versions/master/install.html
2. Clone the TensorFlow repo, run: `git clone https://github.com/tensorflow/tensorflow.git`
3. `$ cd tensorflow`
4. Stay in subdir; first, confirm file's path has not changed: `tensorflow/lite/tools/visualize.py`
5. then run: `$ bazel build tensorflow/lite/tools/visualize` (~1 hour on my 2GHz MacBookPro)
6. then run: `$ bazel-bin/tensorflow/lite/tools/visualize <path-to>/mnist.tflite <path-to>/mnist.html`
(creates the *mnist.html* file at specified <path-to> location)
7. Note: didn't try but maybe can combine steps 5/6, see www.tensorflow.org/lite/devguide

INPUT AND OUTPUT FORMATS IN A TENSOR FLOW LITE MODEL

- ▶ Open the generated mnist.html file
- ▶ inputs: index 16
shape [1, 28, 28, 1]
type FLOAT32
name "x"
- ▶ outputs: index 15
shape [1, 10]
type FLOAT32
name "output"

TensorFlow Lite Model

| | |
|-------------|-----------------|
| filename | mnist.tflite |
| version | 3 |
| description | TOCO Converted. |

Subgraph 0

Inputs/Outputs

| inputs | outputs |
|--------|---------|
| [16] | [15] |

Tensors

| index | name | type | shape | buffer | quantization |
|-------|-----------------------------|---------|-----------------|--------|--|
| 0 | layer_1_conv/Conv2D_bias | FLOAT32 | [6] | 5 | {u'details_type': 0} |
| 1 | layer_1_conv/Relu | FLOAT32 | [1, 28, 28, 6] | 12 | {u'details_type': 0} |
| 2 | layer_1_conv/weight | FLOAT32 | [1, 6, 6, 6] | 13 | {u'details_type': 0} |
| 3 | layer_2_conv/Conv2D_bias | FLOAT32 | [12] | 8 | {u'details_type': 0} |
| 4 | layer_2_conv/Relu | FLOAT32 | [1, 14, 14, 12] | 17 | {u'details_type': 0} |
| 5 | layer_2_conv/weight | FLOAT32 | [12, 5, 5, 6] | 15 | {u'details_type': 0} |
| 6 | layer_3_conv/Conv2D_bias | FLOAT32 | [24] | 10 | {u'details_type': 0} |
| 7 | layer_3_conv/Relu | FLOAT32 | [1, 7, 7, 24] | 9 | {u'details_type': 0} |
| 8 | layer_3_conv/weight | FLOAT32 | [24, 4, 4, 12] | 11 | {u'details_type': 0} |
| 9 | layer_4_fc/MatMul_bias | FLOAT32 | [200] | 3 | {u'details_type': 0} |
| 10 | layer_4_fc/Relu | FLOAT32 | [1, 200] | 7 | {u'details_type': 0} |
| 11 | layer_4_fc/weight/transpose | FLOAT32 | [200, 1176] | 4 | {u'details_type': 0} |
| 12 | layer_5_fc/MatMul_bias | FLOAT32 | [10] | 1 | {u'details_type': 0} |
| 13 | layer_5_fc/add | FLOAT32 | [1, 10] | 6 | {u'details_type': 0} |
| 14 | layer_5_fc/weight/transpose | FLOAT32 | [10, 200] | 2 | {u'details_type': 0} |
| 15 | output | FLOAT32 | [1, 10] | 14 | {u'details_type': 0} |
| 16 | x | FLOAT32 | [1, 28, 28, 1] | 16 | {u'max': [255.0], u'details_type': 0, u'min': [0.0]} |

INPUT AND OUTPUT FORMATS IN A TENSOR FLOW LITE MODEL (2)

- ▶ inputs: index 16
.....
- ▶ outputs: index 15
.....

Subgraph 0

Inputs/Outputs

| inputs | outputs |
|--------|---------|
| [16] | [15] |

Tensors

| index | name |
|-------|------|
|-------|------|

INPUT AND OUTPUT FORMATS IN A TENSOR FLOW LITE MODEL (3)

| | | |
|----|-----------------------------|------------------------|
| 14 | layer_5_fc/weight/transpose | FLOAT32 [10, 200] |
| 15 | output | FLOAT32 [1, 10] |
| 16 | x | FLOAT32 [1, 28, 28, 1] |

► inputs: index 16
shape [1, 28, 28, 1]
type FLOAT32
name "x"

► outputs: index 15
shape [1, 10]
type FLOAT32
name "output"

LABELLING THE INPUT AND OUTPUT IN THE MODEL CODE

► [nex3z/tflite-mnist-android/train.py](#)

[illegible]

INPUT FORMAT FOR MODEL

- ▶ Input shape is [1, 28, 28, 1]
- ▶ The first number is batch size, so pass in a single image for analysis
- ▶ The second and third numbers are the width and height of the image we want to analyze
- ▶ The fourth number is the number of values for each pixel; in this case each pixel will be represented by a single FLOAT32 (specified in type). This will represent a greyscale value from 0.0 (black) to 1.0 (white).
- ▶ (A bit confusingly, there is also the notion of "index" of the input format, index 0 is used since we have just one input format, not to be confused with the batch size of one)

INPUT FORMAT FOR MODEL (2)

```
...
private val imagePixels = IntArray(imageSize)
...

fun classify(
    bitmap: Bitmap,
    success: (Int, Float, Long) -> Unit
) {
    val inputs = FirebaseModelInputs.Builder()
        .add(bitmap.toVector())
        .build()
    val start = System.currentTimeMillis()
    interpreter.run(inputs, options)
        .addOnSuccessListener { outputs -> ... }
        .addOnFailureListener { exception -> throw(exception) }
}

private fun Bitmap.toVector(): Array<Array<Array<FloatArray>>> {
    getPixels(imagePixels, 0, width, 0, 0, width, height)
    return Array(1) {
        Array(imageHeight) { y ->
            Array(imageWidth) { x ->
                floatArrayOf(imagePixels[x + (y * imageWidth)].convertToGreyscale())
            }
        }
    }
}
```

OUTPUT FORMAT FOR MODEL

- ▶ Output shape is [1, 10]
- ▶ The first number is the batch size, conforming with the input format batch
- ▶ The second number represents the number of classifications. Our model will classify the image from 0-9 inclusive, and each classification will have a FLOAT32 (type) value giving the probability that the input image is that digit classification.

```
[7.906156E-16, 2.7768906E-15, 3.609502E-14, 1.0,  
5.1401704E-20, 2.7919297E-12, 7.474837E-17, 6.025316E-17,  
1.3942689E-16, 3.370442E-19]
```

- ▶ In above array, the fourth classification has max value 1.0, so inferred digit is 3
- ▶ (There is also an analogous notion of “index” of the output format, index 0 is used since we also have just one output format.)

OUTPUT FORMAT FOR MODEL (2)

```
fun classify(
    bitmap: Bitmap,
    success: (Int, Float, Long) -> Unit
) {
    val inputs = FirebaseModelInputs.Builder().add(bitmap.toVector()).build()
    val start = System.currentTimeMillis()
    interpreter.run(inputs, options)
        .addOnSuccessListener { outputs -> outputs.map().entries.maxBy { it.value }
            ?.also { maxEntry ->
                success(maxEntry.key, maxEntry.value,
                    System.currentTimeMillis() - start)
            }
        }
        .addOnFailureListener { exception -> throw(exception) }
}
...

private fun FirebaseModelOutputs.map(): Map<Int, Float> {
    return getOutput<Array<FloatArray>>(0)[0]
        .mapIndexed { index, fl -> index to fl }.toMap()
}
```

WHERE IS THE BAD

- ▶ It's all good... but Mark Allison did share one gotcha
 - ▶ "I had got everything working but the accuracy of the digit detection was pretty poor."
- ▶ **The Bad:** Incorrect assumptions about Training Dataset
- ▶ The MNIST images were white digits drawn on black b/g; so subtract original greyscale value calculated from 1.0

```
private fun Int.convertToGreyScale(): Float =  
    1f - ((Color.red(this) + Color.green(this) + Color.blue(this))  
        .toFloat() / 3f / 255f)
```

WRAP UP – ANOTHER EXAMPLE INPUT AND OUTPUT FORMATS

- ▶ In TFLG > Android demo app: camera demo app
 - ▶ If use: quantized Mobilenet TensorFlow Lite model
 - ▶ Input shape is $1 * 224 * 224 * 3$
 - ▶ 1 image in a batch
 - ▶ 224 x 224 width and height of image
 - ▶ 3 bytes (type "uint8" see www.tensorflow.org/lite/tf_ops_compatibility#supported_types)
 - ▶ for the three colors of a pixel
 - ▶ Output shape is $1 * 1001$
 - ▶ the model has 1001 unique categories for the image

APPENDIX 1 – OVERVIEW OF APP LOGIC FLOW

- ▶ In MainActivity, a custom view representing the finger canvas is created. When FingerCanvasView view's onTouchEvent(MotionEvent) is called with ACTION_UP, drawingListener(bitmap) is called
- ▶ In MainActivity, the drawingListener(bitmap) listener function is *defined* to launch a coroutine that calls numberClassifier.classify(bitmap, success); the *success* lambda is *defined*: 1) to take three parameters: the first parameter would be the inferred digit; 2) to launch another coroutine to update the text view at top of app to display that inferred digit passed in
- ▶ The magic happens in the NumberClassifier *classify(bitmap, success)* method. That class has a FirebaseModelInterpreter interpreter. The method first converts the bitmap into FirebaseModelInputs inputs with the expected input format, then calls *interpreter.run(inputs, options)* where options defines the input and output formats; the run call returns a Task<FirebaseModelOutputs> task on which we call addOnSuccessListener(<lambda>) to add a onSuccess lambda that takes a FirebaseModelOutputs outputs parameter. The onSuccess lambda extracts the output array from the FirebaseModelOutputs outputs since we know the output format, then finds the array element (indexed from 0 to 9) with the max value (which is a probability value), then calls the *success* lambda with the index of the max element as the first parameter

APPENDIX 2 – LINKS TO USEFUL BLOG POSTS, TUTORIALS

- ▶ <https://towardsdatascience.com/an-intuitive-guide-to-deep-network-architectures-65fdc477db41>
- ▶ <https://firebase.google.com/docs/ml-kit/android/use-custom-models>